



# Understanding Wide Neural Networks

Jaehoon Lee  
Google Brain

HEP-AI Journal Club  
Feb 5, 2019

# Joint work with



Yasaman Bahri (Brain), Roman Novak (Brain), Jeffrey Pennington (Brain NYC),  
Sam Schoenholz (Brain), Jascha Sohl-Dickstein (Brain), Lechao Xiao (Brain NYC), Greg Yang (MSR)

# Outline



- Motivation
- Deep neural networks as Gaussian processes
  - Formulation / Experiments
- Gradient descent dynamics of wide networks
  - Formulation / Experiments

# Why study wide neural networks?

- Understand effects of overparameterization
- Theoretically simplifying limits (thermodynamic?)
  - Signal propagation
  - Gaussian process correspondence
  - Gradient descent dynamics
- Think in function space ( $f$ ) since parameters ( $w$ ) in a neural network lack direct meaning
  - Random initialization  $p(w)$  induces prior over functions  $p(f)$
  - Wide networks makes function space view more tractable
- Often wide networks perform better

# Is the large width limit uninteresting?

In practice, find that larger width networks trained with stochastic optimization can generalize better.



Generalization gap for five-hidden layer fully-connected networks with variable widths on CIFAR-10.  
Filtered for 100% classification training accuracy.

# Deep neural networks as Gaussian processes

# DEEP NEURAL NETWORKS AS GAUSSIAN PROCESSES

**Jaehoon Lee<sup>\*†</sup>, Yasaman Bahri<sup>\*†</sup>, Roman Novak, Samuel S. Schoenholz,  
Jeffrey Pennington, Jascha Sohl-Dickstein**

Google Brain

{jaehlee, yasamanb, romann, schsam, jpennin, jaschasd}@google.com

- <https://arxiv.org/abs/1711.00165>
- Open source code : <https://github.com/brain-research/nngp>

## Motivations:

- To understand neural networks, can we connect them to objects we better understand?
- An algorithmic aspect: perform Bayesian inference with neural networks?

## Our contributions:

- Correspondence between Gaussian processes and priors for *infinitely wide*, deep neural networks.
- We implement the GP (will refer to as NNGP) and use it to do Bayesian inference. We compare its performance to wide neural networks trained with stochastic optimization on MNIST & CIFAR-10.


# Bayesian treatment of neural networks

- Usual gradient based training of NN : maximum likelihood (or maximum posterior) estimate
- Bayesian deep learning : marginalize over parameter distribution
  - Uncertainty estimates
  - Principled model selection
  - Avoid overfitting (model averaging)
- Why don't we use it then?
  - High computational cost (estimating posterior weight dist)
  - Rely on approximate methods (variational / MCMC)

$$p(w|x, y) = \frac{p(x, y|w)p(w)}{\int p(y|x, w)p(w)dw}$$

# Bayesian treatment of deep neural networks by GPs

- Benefits
  - Uncertainty estimates
  - Principled model selection
  - Avoid overfitting (model averaging)
- Problem
  - High computational cost (estimating posterior weight dist.)
  - Rely on approximate methods (variational / MCMC)
- **Our suggestion**
  - Exact GP equivalence to infinitely wide, deep networks
  - Works for any depth
  - Bayesian inference of NN, without training!



# Reminder: Gaussian Processes


Recall the definition of a Gaussian process:

$z(x) \sim \mathcal{GP}(\mu, K)$ , with mean and covariance functions  $\mu(x), K(x, x')$ , if any finite set of draws,  $[z(x_1), \dots, z(x_n)]^T$ , follows  $\mathcal{N}(\vec{\mu}, \mathbf{K})$  with

$$\vec{\mu} = \begin{bmatrix} \mu(x_1) \\ \vdots \\ \mu(x_n) \end{bmatrix}, \mathbf{K} = \begin{bmatrix} K(x_1, x_1) & \cdots & K(x_1, x_n) \\ \vdots & \ddots & \vdots \\ K(x_n, x_1) & \cdots & K(x_n, x_n) \end{bmatrix}$$


For instance, for the RBF kernel,  $K(x, x') = e^{-\frac{\|x-x'\|^2}{2\sigma^2}}$

Samples from GP with RBF Kernel




# Bayesian inference using a GP prior

Prior with RBF Kernel



Posterior with RBF Kernel



# GP: Bayesian inference

- Bayesian inference involves high-dimensional integration in general.
- For regression, can perform inference exactly because all the integrals are Gaussian

Result (Williams 97) is:

$$\text{Output } z^* | \mathcal{D}, x^* \sim \mathcal{N}(\bar{\mu}, \bar{K})$$

$$\bar{\mu} = K_{x^*, \mathcal{D}} (K_{\mathcal{D}, \mathcal{D}} + \sigma_\epsilon^2 \mathbb{I}_n)^{-1} \mathbf{t}$$

$$\bar{K} = K_{x^*, x^*} - K_{x^*, \mathcal{D}} (K_{\mathcal{D}, \mathcal{D}} + \sigma_\epsilon^2 \mathbb{I}_n)^{-1} K_{x^*, \mathcal{D}}^T$$

Reduces inference to doing linear algebra.

# Shallow Neural Networks and Gaussian Process Priors

*Radford Neal, “Priors for Infinite Networks,” 1994.*

Neal observed that given a neural network (NN) which:

- has a **single hidden layer**
- is **fully-connected**
- has i.i.d. prior over parameters (such that it give a sensible limit)

Then the distribution on its output converges to a Gaussian Process (GP) **in the limit of infinite layer width.**

# Shallow Neural Networks and Gaussian Process Priors

Justification: Central Limit Theorem

$$z_i^1(x) = b_i^1 + \sum_{j=1}^{N_1} W_{ij}^1 x_j^1(x), \quad x_j^1(x) = \phi\left(b_j^0 + \sum_{k=1}^{d_{in}} W_{jk}^0 x_k\right).$$

In the infinite width limit, every finite collection of  $\{z_i^1(x^\mu)\}_{i,\mu}$  will have a joint multivariate Normal distribution: definition of GP.

Let's suppose e.g.:  $W_{i,j}^1 \sim \mathcal{N}(0, \sigma_w^2/N_1)$ ,  $b_i^1 \sim \mathcal{N}(0, \sigma_b^2)$

$$\mu^1(x) = \mathbb{E}[z_i^1(x)] = 0$$

$$K^1(x, x') \equiv \mathbb{E}[z_i^1(x) z_i^1(x')] = \sigma_b^2 + \sigma_w^2 \mathbb{E}[x_i^1(x) x_i^1(x')] \equiv \sigma_b^2 + \sigma_w^2 C(x, x').$$

(Note that outputs are independent because they have Normal joint and zero covariance.)

# Deep Neural Networks and Gaussian Process Priors

What is the prior over functions implied by the prior over parameters, for **deep neural networks**?

Consider a network which:

- is **deep (L layers)**
- is **fully-connected**
- has **i.i.d. prior over parameters (such that it give a sensible limit)**

Then the distribution on its output is also a GP **in the limit of infinite layer width**.

---

$$z_i^l(x) = b_i^l + \sum_{j=1}^{N_l} W_{ij}^l x_j^l(x), \quad x_j^l(x) = \phi(z_j^{l-1}(x)).$$

Suppose (from induction), that  $z_j^{l-1} \sim \mathcal{GP}(0, K^{l-1})$ , and different units  $j$  are independent.

Then similarly, from Central Limit Theorem:  $z_i^l \sim \mathcal{GP}(0, K^l)$

# NNGP covariance function

Recursion relation is:


$$K^l(x, x') = \sigma_b^2 + \sigma_w^2 F_\phi \left( K^{l-1}(x, x'), K^{l-1}(x, x), K^{l-1}(x', x') \right)$$

For some non-linearities, can compute  $F_\phi$  exactly  
(e.g. see Cho and Saul, '09; A. Daniely, et al. '16).

For ReLU:

$$K^l(x, x') = \sigma_b^2 + \frac{\sigma_w^2}{2\pi} \sqrt{K^{l-1}(x, x)K^{l-1}(x', x')} \left( \sin \theta_{x,x'}^{l-1} + (\pi - \theta_{x,x'}^{l-1}) \cos \theta_{x,x'}^{l-1} \right)$$

$$\theta_{x,x'}^l = \cos^{-1} \left( \frac{K^l(x, x')}{\sqrt{K^l(x, x)K^l(x', x')}} \right).$$




ReLU kernel for various depths  
(larger depth gives flatter curves).

# Deep Neural Networks and Gaussian Process Priors


Altogether, for a depth L network, we summarize this:

$$z^L \sim \mathcal{GP}(0, K^L)$$

$$K^L = \sigma_b^2 + \sigma_w^2 F_\phi(K^{L-1})$$



Samples from a GP neural network prior with depth 10.



# Reference for more formal treatment


- A. Matthews et al., ICLR 2018
  - Gaussian Process Behaviour in Wide Deep Neural Networks
  - <https://arxiv.org/abs/1804.11271>
- R. Novak et al., ICLR 2019
  - Bayesian Deep Convolutional Networks with Many Channels are Gaussian Processes
  - <https://arxiv.org/abs/1810.05148>
  - Appendix E

# Experiments

# Experimental setup

- Datasets: MNIST, CIFAR-10
- Permutation invariant, fully-connected model, ReLU/Tanh activation function
- Trained on mean squared loss
- Targets are one-hot encoded, zero-mean and treated as regression target
  - incorrect class -0.1, correct class 0.9
- Hyperparameter optimized using random / grid search
  - Weight / bias variances, optimization hyperparameters (for NN)
- NN: ‘SGD’ trained opposed to Bayesian training. In practice, Adam optimizer was used (qualitatively similar).
- NNGP: standard exact Gaussian process regression, 10 independent outputs

# Performance of wide networks approaches NNGP



Accuracy of finite-width, fully-connected deep NN + SGD →  
NNGP with exact Bayesian inference









# Finite width networks trained with SGD vs NNGP

Num training	Model (ReLU)	Test accuracy	Model (tanh)	Test accuracy
MNIST:1k	NN-2-5000-3.19-0.00	0.9252	NN-2-1000-0.60-0.00	0.9254
	GP-20-1.45-0.28	<b>0.9279</b>	GP-20-1.96-0.62	0.9266
MNIST:10k	NN-2-2000-0.42-0.16	0.9771	NN-2-2000-2.41-1.84	0.9745
	GP-7-0.61-0.07	0.9765	GP-2-1.62-0.28	<b>0.9773</b>
MNIST:50k	NN-2-2000-0.60-0.44	0.9864	NN-2-5000-0.28-0.34	0.9857
	GP-1-0.10-0.48	0.9875	GP-1-1.28-0.00	<b>0.9879</b>
CIFAR:1k	NN-5-500-1.29-0.28	0.3225	NN-1-200-1.45-0.12	0.3378
	GP-7-1.28-0.00	0.3608	GP-50-2.97-0.97	<b>0.3702</b>
CIFAR:10k	NN-5-2000-1.60-1.07	0.4545	NN-1-500-1.48-1.59	0.4429
	GP-5-2.97-0.28	<b>0.4780</b>	GP-7-3.48-2.00	0.4766
CIFAR:45k	NN-3-5000-0.53-0.01	0.5313	NN-2-2000-1.05-2.08	0.5034
	GP-3-3.31-1.86	<b>0.5566</b>	GP-3-3.48-1.52	0.5558

$$\text{NN-depth-width-}\sigma_w^2-\sigma_b^2 \quad \text{GP-depth-}\sigma_w^2-\sigma_b^2$$

# NNGP hyperparameter dependence

Test accuracy





# Uncertainty

- Neural networks are good at making predictions, but does not naturally provide uncertainty estimates
- Bayesian methods incorporates uncertainty
- In domains where uncertainty of prediction is important, GP has been useful
- In NNGP, uncertainty of NN's prediction is captured by variance in output

$$\bar{K} = K_{x^*, x^*} - K_{x^*, \mathcal{D}}(K_{\mathcal{D}, \mathcal{D}} + \sigma_\epsilon^2 \mathbb{I}_n)^{-1} K_{x^*, \mathcal{D}}^T$$

# Uncertainty: how good are the estimates?



X: predicted uncertainty

Y: realized MSE

\* averaged over 100  
points binned by  
predicted uncertainty

Empirical error is well correlated with uncertainty predictions

# Log marginal likelihood (model selection)

$$\log p(t|\theta) = -\frac{1}{2}t^T(K_{DD}(\theta) + \sigma_\epsilon^2 \mathbb{I})^{-1}t - \frac{1}{2}\log \det(K_{DD}(\theta) + \sigma_\epsilon^2 \mathbb{I}) + \text{const}$$


- Neural network hyperparameters: depth, weight / bias variance, non-linearity
- No validation set is required to select model hyperparameters. Evaluate on train data.
- $K_{DD}$  is deterministic and differentiable, implemented in Tensorflow. Can backprop!

# Future works

**NNGP correspondence opens up interesting angles to further analyze deep neural networks.**

Published as a conference paper at ICLR 2019

BAYESIAN DEEP CONVOLUTIONAL NETWORKS WITH  
MANY CHANNELS ARE GAUSSIAN PROCESSES

Roman Novak <sup>†</sup>, Lechao Xiao <sup>† \*</sup>, Jaehoon Lee <sup>‡ \*</sup>, Yasaman Bahri <sup>† \*</sup>, Greg Yang <sup>○</sup>,  
Daniel A. Abolafia, Jeffrey Pennington, Jascha Sohl-Dickstein

Google Brain, <sup>○</sup>Microsoft Research AI

{romann, xlc, jaehlee, yasamanb, gregyang@microsoft.com,  
danabo, jpennin, jaschasd}@google.com

- Practical usage of NNGP
- Extension to other network architectures
  - **Convolutional / Residual** [Novak et al., ICLR 2019, Garriga-Alonso et al., ICLR 2019]
  - Batch normalization, self-attention, recurrent, ...
- Systematic finite width correction

# Gradient descent dynamics of wide networks

---

# Gaussian Predictions from Gradient Descent Training of Wide Neural Networks

---

NeurIPS Bayesian Deep  
Learning Workshop 2019

Jaehoon Lee\*, Lechao Xiao\*, Jascha Sohl-Dickstein, Jeffrey Pennington  
Google Brain  
{jaehlee, xlc, jaschasd, jpennin}@google.com

---


**Wide neural networks of any depth  
evolve as linear models under gradient descent**

---

Available at arXiv soon

Jaehoon Lee<sup>\*1</sup> Lechao Xiao<sup>\*1</sup> Sam Schoenholz<sup>1</sup>  
Yasaman Bahri<sup>1</sup> Jascha Sohl-Dickstein<sup>1</sup> Jeffrey Pennington<sup>1</sup>

# Recall : empirical observations



Accuracy of finite-width, fully-connected deep NN + SGD →  
NNGP with exact Bayesian inference

How similar is gradient descent based training to the Bayesian inference?

## Motivations:

- Bayesian inference VS gradient descent training
- Tractable learning dynamics of deep neural networks

## Our contributions:

- Wide neural networks' training dynamics under gradient descent become surprisingly simple
  - Effectively replace NN by its first-order Taylor expansion around init parameters
  - Linear model captures the NN training dynamics
- Analytic dynamics for MSE loss, simple generalization to xent loss / momentum optimizer / practical networks (wide residual network)
- Analytic output distribution dynamics for MSE loss: not equal to NNGP posterior

# Gradient descent dynamics (continuous time)

$$\mathcal{L} = \sum_{(x,y) \in \mathcal{D}} \ell(f_t(x, \theta), y).$$

$$\dot{\theta}_t = -\eta \nabla_{\theta} f_t(\mathcal{X})^T \nabla_{f_t(\mathcal{X})} \mathcal{L}$$

$$\dot{f}_t(\mathcal{X}) = \nabla_{\theta} f_t(\mathcal{X}) \dot{\theta}_t = -\eta \hat{\Theta}_t(\mathcal{X}, \mathcal{X}) \nabla_{f_t(\mathcal{X})} \mathcal{L}$$

$$\hat{\Theta}_t = \nabla_{\theta} f_t(\mathcal{X}) \nabla_{\theta} f_t(\mathcal{X})^T = \sum_{l=1}^{L+1} \nabla_{\theta^l} f_t(\mathcal{X}) \nabla_{\theta^l} f_t(\mathcal{X})^T.$$

Neural Tangent  
Kernel (NTK)  
[Jacot et al. 2018]

# Linearized networks

$$f_t^{\text{lin}}(x) \equiv f_0(x) + \nabla_{\theta} f_0(x) \omega_t$$

$$\omega_t \equiv \theta_t - \theta_0$$

$$\dot{\omega}_t = -\eta \nabla_{\theta} f_0(\mathcal{X})^T \nabla_{f_t^{\text{lin}}(\mathcal{X})} \mathcal{L}$$

$$\dot{f}_t^{\text{lin}}(x) = -\eta \hat{\Theta}_0(x, \mathcal{X}) \nabla_{f_t^{\text{lin}}(\mathcal{X})} \mathcal{L}.$$

Dynamics fully determined by initialization objects: **simple ODE**

# Tractable dynamics for wide networks

- Remarkably Jacot et al. 2018 showed that

$$\sup_{t \in [0, T]} \|\hat{\Theta}_t - \hat{\Theta}_0\|_F = O(\min\{n_1, \dots, n_L\}^{-1/2})$$

- For MSE loss, we also show that

$$\sup_{t \in [0, T]} \|f_t(\mathcal{X}) - f_t^{\text{lin}}(\mathcal{X})\|_2 = O(\sup_{t \in [0, T]} \|\hat{\Theta}_t - \hat{\Theta}_0\|_F),$$

- Linearized networks training dynamics converges to that of original network as width increases

# Predictive output distribution



- Sample-then-optimize posterior sampling (Matthews et al., 2017)
  - Randomly initialize networks
  - Optimize (via GD) using training data
  - Predictive output distribution over ensemble of different initialization
- For wide networks
  - Only optimize readout weights : interpolation between prior and posterior of NNGP
  - Optimize all the weights: As width increases, **ensembles** of random wide neural networks trained with (stochastic) gradient descent converges to a Gaussian process

$$\mu(x) = \Theta(x, \mathcal{X})\Theta^{-1}(I - e^{-\eta\Theta t})\mathcal{Y} \quad (15)$$

$$\begin{aligned} \Sigma(x) &= \mathcal{K}(x, x) - 2\Theta(x, \mathcal{X})\Theta^{-1}(I - e^{-\eta\Theta t})\mathcal{K}(x, \mathcal{X})^T \\ &\quad + \Theta(x, \mathcal{X})\Theta^{-1}(I - e^{-\eta\Theta t})\mathcal{K}\Theta^{-1}(I - e^{-\eta\Theta t})\Theta(x, \mathcal{X})^T. \end{aligned} \quad (16)$$


# Experiments

# NN posterior vs GP posterior




# Comparison of training dynamics linearized network vs original network

**FC / MSE / GD**



**WResNet\* / xent / momentum**



CIFAR binary classification with 128 samples

**Thank you! Questions?**

# NTK parameterization of NN

Conventional

$$z^l = W^l x^l + b^l$$

$$W_{ij}^l \sim \mathcal{N}(0, \sigma_w^2 / n^l)$$

NTK [Jacot et al 2018]

$$z^l = \frac{1}{\sqrt{n^l}} \tilde{W}^l x^l + b^l$$

$$\tilde{W}_{ij}^l \sim \mathcal{N}(0, \sigma_w^2)$$

Computes the same functions / modifies dynamics / universal learning rates (absorb 1/n)

# Deep Neural Networks and Gaussian Process Priors

$$z_i^l(x) = b_i^l + \sum_{j=1}^{N_l} W_{ij}^l x_j^l(x), \quad x_j^l(x) = \phi(z_j^{l-1}(x)).$$

$$K^l(x, x') \equiv \mathbb{E} [z_i^l(x) z_i^l(x')] = \sigma_b^2 + \sigma_w^2 \mathbb{E}_{z_i^{l-1} \sim \mathcal{GP}(0, K^{l-1})} [\phi(z_i^{l-1}(x)) \phi(z_i^{l-1}(x'))]$$

The calculation of the expectation is a 2D Gaussian integral:

$$K^l(x, x') = \sigma_b^2 + \sigma_w^2 \mathcal{Z}^{-1} \int du_1 du_2 \phi(u_1) \phi(u_2) \exp \left( -\frac{1}{2} [u_1, u_2] \begin{bmatrix} K^{l-1}(x, x) & K^{l-1}(x, x') \\ K^{l-1}(x, x') & K^{l-1}(x', x') \end{bmatrix}^{-1} \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} \right)$$

As a result:

$$K^l(x, x') = \sigma_b^2 + \sigma_w^2 F_\phi \left( K^{l-1}(x, x'), K^{l-1}(x, x), K^{l-1}(x', x') \right)$$

Base case in the recursion:

$$K^0(x, x') = \mathbb{E} [z_j^0(x) z_j^0(x')] = \sigma_b^2 + \sigma_w^2 \left( \frac{x \cdot x'}{d_{\text{in}}} \right)$$